

HBASE – HIVE INTEGRATION

By.Mr.Gopal Krishna
Sr.Hadoop Technical Architect
CCA 175-Spark and Hadoop Developer Certified Consultant

+++++

FROM HBASE SIDE

+++++

```
hbase(main):002:0> create 'hbasehive_emptable','empdetails'
```

```
0 row(s) in 1.9190 seconds
```

```
=> Hbase::Table - hbasehive_emptable
```

```
hbase(main):003:0>
```

```
hbase(main):003:0> scan 'hbasehive_emptable'
```

```
ROW          COLUMN+CELL
```

```
0 row(s) in 0.1550 seconds
```

```
hbase(main):004:0>
```

+++++

FROM HIVE SIDE

+++++

```
hive> create database hivehbasedb;
```

```
OK
```

```
Time taken: 0.493 seconds
```

```
hive> use hivehbasedb;
```

```
OK
```

```
Time taken: 0.02 seconds
```

ADD THE BELOW THREE JAR FILES – FOR HBASE –HIVE INTEGRATION

```
hive> add jar ${env:HIVE_HOME}/lib/zookeeper-3.4.6.jar;
```

```
Added [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/zookeeper-3.4.6.jar] to class path
```

```
Added resources: [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/zookeeper-3.4.6.jar]
```

```
hive> add jar ${env:HIVE_HOME}/lib/guava-14.0.1.jar;
```

```
Added [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/guava-14.0.1.jar] to class path
```

```
Added resources: [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/guava-14.0.1.jar]
```

```
hive> add jar ${env:HIVE_HOME}/lib/hive-hbase-handler-1.2.1.jar;
```

```
Added [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/hive-hbase-handler-1.2.1.jar] to class path
```

```
Added resources: [/home/gopalkrishna/INSTALL/apache-hive-1.2.1-bin/lib/hive-hbase-handler-1.2.1.jar]
```

```
hive>
```

CREATE A HIVE EXTERNAL TABLE WITH HBASE COLUMNS MAPPING

```
hive> CREATE EXTERNAL TABLE hive_emptable(key int,  
      firstname string,  
      lastname string,  
      yrsofexp int,  
      desg string)  
  > STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'  
  > WITH SERDEPROPERTIES  
  > (  
  >   "hbase.columns.mapping"=  
":key,empdetails:firstname,empdetails:lastname,empdetails:yrsofexp,empdetails:de  
sg")  
  > TBLPROPERTIES ("hbase.table.name"="hbasehive_emptable");
```

OK

Time taken: 0.464 seconds

```
hive> describe formatted hive_emptable;
```

OK

#	col_name	data_type	comment
	key	int	from deserializer
	firstname	string	from deserializer
	lastname	string	from deserializer
	yrsofexp	int	from deserializer
	desg	string	from deserializer

Detailed Table Information

Database: hivehbasedb
Owner: gopalkrishna
CreateTime: Thu Sep 14 00:46:48 PDT 2017
LastAccessTime: UNKNOWN
Protect Mode: None
Retention: 0
Location: hdfs://localhost:8020/user/hive/warehouse/hivehbasedb.db/hive_emptable
Table Type: EXTERNAL_TABLE
Table Parameters:

COLUMN_STATS_ACCURATE	false
EXTERNAL	TRUE
hbase.table.name	hbasehive_emptable
numFiles	0
numRows	-1
rawDataSize	-1
storage_handler	org.apache.hadoop.hive.hbase.HBaseStorageHandler
totalSize	0
transient_lastDdlTime	1505375208

Storage Information

SerDe Library: org.apache.hadoop.hive.hbase.HBaseSerDe
InputFormat: null
OutputFormat: null
Compressed: No
Num Buckets: -1
Bucket Columns: []
Sort Columns: []
Storage Desc Params:

```
hbase.columns.mapping
:key,empdetails:firstname,empdetails:lastname,empdetails:yrsofexp,empdetails:desg
serialization.format      1
```

Time taken: 0.187 seconds, Fetched: 39 row(s)

hive>

BEFORE LOADING THE DATA INTO HIVE TABLE

hive> select * from hive_emptable;

OK

Time taken: 1.301 seconds

hive>

TO LOAD THE DATA INTO HIVE TABLE

hive> LOAD DATA LOCAL INPATH

'/home/gopalkrishna/PRAC/HBASE/HBase_Hive_InputData.log' INTO TABLE
hive_emptable;

FAILED: SemanticException [Error 10101]: A non-native table cannot be used as target for LOAD

hive>

hive> CREATE EXTERNAL TABLE initable(empid int,

- > firstname string,
- > lastname string,
- > yrsofexp int,
- > desg string)
- > ROW FORMAT DELIMITED
- > FIELDS TERMINATED BY '\t'
- > LINES TERMINATED BY '\n'
- > STORED AS TEXTFILE;

OK

Time taken: 0.105 seconds

hive> LOAD DATA LOCAL INPATH

'/home/gopalkrishna/PRAC/HBASE/HBase_Hive_InputData.log' INTO TABLE initable;

Loading data to table hivehbasedb.initable

Table hivehbasedb.initable stats: [numFiles=2, numRows=0, totalSize=1142, rawDataSize=0]

OK

Time taken: 0.822 seconds

hive> SELECT * FROM initable;

OK

1000	RAVI	VARMA	9	TECHNOLOGY LEAD
1001	KOUSIK	MISTRA	11	ENTRYLEVEL ARCHITECT
1002	SAIED	FAISAL	21	DELIVERY HEAD
1003	SOURAVJAIN		18	SENIOUR PROJECT MANAGER
1004	CHRIS	WAYNE	15	TEST MANAGER
1005	ANEES	FLEXWALA	19	CONFIGURATION CONTROLLER
1006	GOPAL	KRISHNA	13	SR.BIGDATA ARCHITECT
1007	RAMYA	SRI	8	BIGDATA LEAD
1008	MUNI	KRISHNA	5	SR.HADOOP DEVELOPER
1009	RAMYA	SRI	7	SR.SPARK DEVELOPER
1010	RAM	CHANDAR	16	PROJECT MANAGER
1011	AKHILA	PRIYA	3	HADOOP DEVELOPER
1012	RAHUL	DEV	6	HADOOP ADMINISTRATOR
1013	AMIT	MISRA	4	BIGDATA TESTER
1014	PUJA	YERUNKAR	11	BIGDATA TEST LEAD
1000	RAVI	VARMA	9	TECHNOLOGY LEAD

1001	KOUSIK	MISTRA	11	ENTRYLEVEL ARCHITECT
1002	SAIED	FAISAL	21	DELIVERY HEAD
1003	SOURAVJAIN		18	SENIOUR PROJECT MANAGER
1004	CHRIS	WAYNE	15	TEST MANAGER
1005	ANEES	FLEXWALA	19	CONFIGURATION CONTROLLER
1006	GOPAL	KRISHNA	13	SR.BIGDATA ARCHITECT
1007	RAMYA	SRI	8	BIGDATA LEAD
1008	MUNI	KRISHNA	5	SR.HADOOP DEVELOPER
1009	RAMYA	SRI	7	SR.SPARK DEVELOPER
1010	RAM	CHANDAR	16	PROJECT MANAGER
1011	AKHILA	PRIYA	3	HADOOP DEVELOPER
1012	RAHUL	DEV	6	HADOOP ADMINISTRATOR
1013	AMIT	MISRA	4	BIGDATA TESTER
1014	PUJA	YERUNKAR	11	BIGDATA TEST LEAD

Time taken: 0.127 seconds, Fetched: 30 row(s)

TO LOAD THE DATA INTO hive_emptable(hbase_mapping table) FROM initable(Non_Mapping Table)

hive> INSERT OVERWRITE TABLE hive_emptable SELECT * FROM initable;

Query ID = gopalkrishna_20170914005226_af713c14-1e4d-45c2-9cec-2819fce28245

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1505372852042_0001, Tracking URL =

http://ubuntu:8088/proxy/application_1505372852042_0001/

Kill Command = /home/gopalkrishna/INSTALL/hadoop-2.6.0/bin/hadoop job -kill

job_1505372852042_0001

Hadoop job information for Stage-0: number of mappers: 1; number of reducers: 0

2017-09-14 00:53:24,750 Stage-0 map = 0%, reduce = 0%

2017-09-14 00:54:01,402 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 8.02 sec

MapReduce Total cumulative CPU time: 8 seconds 20 msec

Ended Job = job_1505372852042_0001

MapReduce Jobs Launched:

Stage-Stage-0: Map: 1 Cumulative CPU: 8.41 sec HDFS Read: 5653 HDFS Write: 0 SUCCESS

Total MapReduce CPU Time Spent: 8 seconds 410 msec

OK

Time taken: 97.971 seconds

hive>

AFTER LOADING THE DATA INTO "hive_emptable"

hive> SELECT * FROM hive_emptable;

OK

1000	RAVI	VARMA	9	TECHNOLOGY LEAD
1001	KOUSIK	MISTRA	11	ENTRYLEVEL ARCHITECT
1002	SAIED	FAISAL	21	DELIVERY HEAD
1003	SOURAVJAIN		18	SENIOUR PROJECT MANAGER
1004	CHRIS	WAYNE	15	TEST MANAGER
1005	ANEES	FLEXWALA	19	CONFIGURATION CONTROLLER
1006	GOPAL	KRISHNA	13	SR.BIGDATA ARCHITECT
1007	RAMYA	SRI	8	BIGDATA LEAD
1008	MUNI	KRISHNA	5	SR.HADOOP DEVELOPER
1009	RAMYA	SRI	7	SR.SPARK DEVELOPER
1010	RAM	CHANDAR	16	PROJECT MANAGER
1011	AKHILA	PRIYA	3	HADOOP DEVELOPER
1012	RAHUL	DEV	6	HADOOP ADMINISTRATOR
1013	AMIT	MISRA	4	BIGDATA TESTER
1014	PUJA	YERUNKAR	11	BIGDATA TEST LEAD

Time taken: 0.297 seconds, Fetched: 15 row(s)

hive>

+++++

AFTER DATA LOAD IN HIVE TABLE - VERIFICATION FROM HBASE SIDE

+++++

GO TO HBASE SHELL AND TRIGGER THE scan COMMAND LIKE BELOW

hbase(main):005:0> scan 'hbasehive_emptable'

ROW

COLUMN+CELL

1000	column=empdetails:desg, timestamp=1505375640630, value=TECHNOLOGY LEAD
1000	column=empdetails:firstname, timestamp=1505375640630, value=RAVI
1000	column=empdetails:lastname, timestamp=1505375640630, value=VARMA
1000	column=empdetails:yrsofexp, timestamp=1505375640630, value=9
1001	column=empdetails:desg, timestamp=1505375640630, value=ENTRYLEVEL ARCHITECT
1001	column=empdetails:firstname, timestamp=1505375640630, value=KOUSIK
1001	column=empdetails:lastname, timestamp=1505375640630, value=MISTRA
1001	column=empdetails:yrsofexp, timestamp=1505375640630, value=11
1002	column=empdetails:desg, timestamp=1505375640630, value=DELIVERY HEAD
1002	column=empdetails:firstname, timestamp=1505375640630, value=SAIED
1002	column=empdetails:lastname, timestamp=1505375640630, value=FAISAL
1002	column=empdetails:yrsofexp, timestamp=1505375640630, value=21
1003	column=empdetails:desg, timestamp=1505375640630, value=SENIOUR PROJECT MANAGER
1003	column=empdetails:firstname, timestamp=1505375640630, value=SOURAV
1003	column=empdetails:lastname, timestamp=1505375640630, value=JAIN
1003	column=empdetails:yrsofexp, timestamp=1505375640630, value=18
1004	column=empdetails:desg, timestamp=1505375640630, value=TEST MANAGER
1004	column=empdetails:firstname, timestamp=1505375640630, value=CHRIS
1004	column=empdetails:lastname, timestamp=1505375640630, value=WAYNE
1004	column=empdetails:yrsofexp, timestamp=1505375640630, value=15
1005	column=empdetails:desg, timestamp=1505375640630, value=CONFIGURATION CONTROLLER
1005	column=empdetails:firstname, timestamp=1505375640630, value=ANEES
1005	column=empdetails:lastname, timestamp=1505375640630, value=FLEXWALA
1005	column=empdetails:yrsofexp, timestamp=1505375640630, value=19
1006	column=empdetails:desg, timestamp=1505375640630, value=SR.BIGDATA ARCHITECT
1006	column=empdetails:firstname, timestamp=1505375640630, value=GOPAL
1006	column=empdetails:lastname, timestamp=1505375640630, value=KRISHNA
1006	column=empdetails:yrsofexp, timestamp=1505375640630, value=13
1007	column=empdetails:desg, timestamp=1505375640630, value=BIGDATA LEAD
1007	column=empdetails:firstname, timestamp=1505375640630, value=RAMYA
1007	column=empdetails:lastname, timestamp=1505375640630, value=SRI
1007	column=empdetails:yrsofexp, timestamp=1505375640630, value=8
1008	column=empdetails:desg, timestamp=1505375640630, value=SR.HADOOP DEVELOPER
1008	column=empdetails:firstname, timestamp=1505375640630, value=MUNI
1008	column=empdetails:lastname, timestamp=1505375640630, value=KRISHNA
1008	column=empdetails:yrsofexp, timestamp=1505375640630, value=5
1009	column=empdetails:desg, timestamp=1505375640630, value=SR.SPARK DEVELOPER
1009	column=empdetails:firstname, timestamp=1505375640630, value=RAMYA
1009	column=empdetails:lastname, timestamp=1505375640630, value=SRI
1009	column=empdetails:yrsofexp, timestamp=1505375640630, value=7
1010	column=empdetails:desg, timestamp=1505375640630, value=PROJECT MANAGER
1010	column=empdetails:firstname, timestamp=1505375640630, value=RAM
1010	column=empdetails:lastname, timestamp=1505375640630, value=CHANDAR
1010	column=empdetails:yrsofexp, timestamp=1505375640630, value=16
1011	column=empdetails:desg, timestamp=1505375640630, value=HADOOP DEVELOPER
1011	column=empdetails:firstname, timestamp=1505375640630, value=AKHILA
1011	column=empdetails:lastname, timestamp=1505375640630, value=PRIYA
1011	column=empdetails:yrsofexp, timestamp=1505375640630, value=3
1012	column=empdetails:desg, timestamp=1505375640630, value=HADOOP ADMINISTRATOR
1012	column=empdetails:firstname, timestamp=1505375640630, value=RAHUL

```

1012      column=empdetails:lastname, timestamp=1505375640630, value=DEV
1012      column=empdetails:yrsofexp, timestamp=1505375640630, value=6
1013      column=empdetails:desg, timestamp=1505375640630, value=BIGDATA TESTER
1013      column=empdetails:firstname, timestamp=1505375640630, value=AMIT
1013      column=empdetails:lastname, timestamp=1505375640630, value=MISRA
1013      column=empdetails:yrsofexp, timestamp=1505375640630, value=4
1014      column=empdetails:desg, timestamp=1505375640630, value=BIGDATA TEST LEAD
1014      column=empdetails:firstname, timestamp=1505375640630, value=PUJA
1014      column=empdetails:lastname, timestamp=1505375640630, value=YERUNKAR
1014      column=empdetails:yrsofexp, timestamp=1505375640630, value=11
15 row(s) in 0.6630 seconds

```

hbase(main):006:0>

INSERT 1 MORE RECORD FROM HBASE SIDE => VERIFY THE SAME RECORD IS REFLECTED IN HIVE TABLE (hive_emptable) OR NOT

hbase(main):006:0> put 'hbasehive_emptable','1100','empdetails:firstname','GOPAL'

0 row(s) in 0.2460 seconds

hbase(main):007:0> put 'hbasehive_emptable','1100','empdetails:lastname','KRISHNA'

0 row(s) in 0.0320 seconds

hbase(main):008:0> put 'hbasehive_emptable','1100','empdetails:yrsofexp',13

0 row(s) in 0.0200 seconds

hbase(main):009:0> put 'hbasehive_emptable','1100','empdetails:desg','SR.HADOOP ARCHITECT'

0 row(s) in 0.0080 seconds

hbase(main):010:0>

HIVE SIDE VERIFICATION

hive> SELECT * FROM hive_emptable;

```

OK
1000  RAVI    VARMA  9      TECHNOLOGY LEAD
1001  KOUSIK  MISTRA 11     ENTRYLEVEL ARCHITECT
1002  SAIED   FAISAL 21     DELIVERY HEAD
1003  SOURAVJAIN 18    SENIOUR PROJECT MANAGER
1004  CHRIS   WAYNE  15     TEST MANAGER
1005  ANEES   FLEXWALA 19     CONFIGURATION CONTROLLER
1006  GOPAL   KRISHNA 13     SR.BIGDATA ARCHITECT
1007  RAMYA   SRI     8      BIGDATA LEAD
1008  MUNI    KRISHNA 5       SR.HADOOP DEVELOPER
1009  RAMYA   SRI     7      SR.SPARK DEVELOPER
1010  RAM     CHANDAR 16     PROJECT MANAGER
1011  AKHILA  PRIYA  3      HADOOP DEVELOPER
1012  RAHUL   DEV     6      HADOOP ADMINISTRATOR
1013  AMIT    MISRA   4      BIGDATA TESTER
1014  PUJA    YERUNKAR 11     BIGDATA TEST LEAD
1100  GOPAL   KRISHNA 13     SR.HADOOP ARCHITECT

```

Time taken: 0.341 seconds, Fetched: 16 row(s)

hive>